

Using the Computer Science Cluster

Alan Crawford
Lars Nesheim

June 27, 2018

Computer Science (CS) High-Performance Computing (HPC) Cluster

- CS cluster has over 5200 nodes.
- Designed to run large scale computing jobs in **batch mode** (non-interactive mode).
 - **Limited graphics**-based interactive computing services.
- Users should set up their code to run in batch mode.
- Users who primarily need graphics-based interactive computing should use Economics Department cluster, ISD services, or desktop computers.

More information

- CS cluster webpage: http://hpc.cs.ucl.ac.uk/how_to_login/.
 - Get username and password from IT support,
- Econ GitHub page: <https://github.com/UCL/ECON-CLUSTER>
 - Contact Fatima or Andrew to be added to the UCL user group.
- If you need help,
 - CS support: cluster-support@cs.ucl.ac.uk
 - ECON IT support: economics.it@ucl.ac.uk

Obtain account

- Fill out and submit application form (see Fatima or Andrew).
- Wait 5-7 days and go to CS Cluster office: (4.20 Malet Place Engineering Building).
- They will set up two accounts
 - 1 CS departmental account (for remote access).
 - 2 CS Cluster account (to use cluster)
- By default, both accounts have same username and password.

Log on to cluster headnode: Thinlinc, ssh, or sftp

Log on to compute node: interactive session

Submit batch job to compute node: non-interactive session

Submit batch job to compute node: non-interactive session

Job status

Disk storage

Software available

Memory management

Graphics on jake or elwood

Cluster rules

- Never run jobs on head nodes (wise, vic, tails)
- Store important data and/or large data on /SAN/economics/...
- Avoid high frequency reading/writing of large datafiles to/from disk
- If you need support, send email to cluster support, go visit them in person.
- Update wiki page with tips for colleagues about how to solve common problems.
-